

Global Namespaces Scale as Well as Decoupled Namespaces

Paper ID: 165

Total Pages: 12

Abstract

Today's large and highly-parallel workloads are bottlenecked by metadata services because many processes end up accessing the same shared resource. In HPC, state-of-the-art file systems are abandoning POSIX because the synchronization and serialization overheads are too costly – and sometimes even unnecessary – for some of their applications. While the performance benefits are plain for these users, other applications that rely on stronger consistency must be re-written or deployed on a different system. We present Cudele, a programmable file system that supports different degrees of consistency and fault tolerance within the same namespace. First, we take a POSIX compliant file system and relax the consistency constraints by implementing two metadata designs: delayed metadata update merge into the global namespace and client local views of metadata updates. Second, we store the consistency and fault tolerance semantics in inodes so that subtrees within the same namespace can be optimized for different workloads; the inodes are programmable so that clients understand how to access the metadata in a subtree. Third, we present a theoretical framework for a metadata service and a working implementation that provides the lowest common denominator needed to implement a wide range of consistency/fault tolerance semantics that can be benchmarked all on the same system.

1. Introduction

Today's client-server based file system metadata services have scalability problems. It takes a lot of resources to server POSIX metadata requests and applications perform better with dedicated metadata servers [6, 8]. This is fine for small workloads and file systems but as the system scales provisioning a metadata server for every client is expensive and complicated.




Figure 1. Administrators can assign consistency and fault tolerance policies to subtrees to get the benefits of some of the state-of-the-art HPC architectures.

Current hardware evolution and the rise of software-defined storage, which uses techniques like erasure coding, replication, and partitioning, have ushered a new era of HPC computing; architectures are transitioning from complex storage stacks with burst buffer, file system, object store, and tape tiers to a two layer stack with just a burst buffer and object store [4]. This trend exacerbates the metadata scalability problem and has given rise to the serverless metadata services.

HPC workloads are so metadata intensive that new management techniques are advocating reducing synchronization and serialization overheads by transferring these responsibilities to the client. We call this approach decoupling the namespace and the semantics of consistency differ amongst systems.

We propose subtree policies, an interface that lets future programmers control how the system manages different parts of the namespace. For performance one subtree can adopt weaker consistency semantics while another subtree can retain the rigidity of POSIX's strong consistency. Figure 1 shows an example setup where a single global namespace has directories for applications designed for different, state-of-the-art HPC architectures. Our system supports 3 forms of consistency and 2 forms of fault tolerance giving the administrator a wide range of policies and optimizations depending on the application's needs.

2. Related Work

The bottlenecks associated with accessing POSIX file system metadata are not limited to HPC workloads and the same challenges that plagued these systems for years are finding their way into the cloud. Workloads that deal with many small files (*e.g.*, log processing and database queries [9]) and large numbers of simultaneous clients (*e.g.*, MapReduce jobs [5]), are subject to the scalability of the metadata service. The biggest challenge is that whenever a file is touched the client must access the file’s metadata and maintaining a file system namespace imposes small, frequent accesses on the underlying storage system [7]. Unfortunately, scaling file system metadata is a well-known problem and solutions for scaling data IO do not work for metadata IO [1–3, 7, 10]. There are two approaches for improving the performance of metadata access.

2.1 Metadata Load Balancing

One approach for improving metadata performance and scalability is to alleviate overloaded servers by load balancing metadata IO across a cluster. Common techniques include partitioning metadata when there are many writes and replicating metadata when there are many reads. For example, IndexFS partitions directories and clients write to different partitions by grabbing leases and caching ancestor metadata for path traversal; it does well for strong scaling because servers can keep more inodes in the cache which results in less RPCs. Alternatively, ShardFS replicates directory state so servers do not need to contact peers for path traversal; it does well for read workloads because all file operations only require 1 RPC and for weak scaling because requests will never incur extra RPCs due to a full cache. CephFS employs both techniques to a lesser extent; directories can be replicated or sharded but the caching and replication policies do not change depending on the balancing technique. Despite the performance benefits these techniques add complexity and jeopardize the robustness and performance characteristics of the metadata service because the systems now need (1) policies to guide the migration decisions and (2) mechanisms to address inconsistent states across servers.

Setting policies for migrations is arguably more difficult than adding the migration mechanisms themselves. For example, IndexFS and CephFS use the GIGA+ [?] technique for partitioning directories at a predefined threshold and using lazy synchronization to redirect queries to the server that “owns” the targeted metadata. Determining when to partition directories and when to migrate the directory fragments are policies that vary between systems: GIGA+ partitions directories when the size reaches a certain number of files and migrates directory fragments immediately; CephFS partitions directories when they reach a threshold size or when the write temperature reaches a certain value and migrates directory fragments when the hosting server has more load than the other servers in the metadata cluster. Another policy

is when and how to replicate directory state; ShardFS replicates immediately and pessimistically while CephFS replicates only when the read temperature reaches a threshold. There is a wide range of policies and it is difficult to address with tunables and hard-coded design decisions.

In addition to the policies, distributing metadata across a cluster requires distributed transactions and cache coherence protocols to ensure strong consistency (*e.g.*, POSIX). For example, ShardFS pessimistically replicates directory state and uses optimistic concurrency control for conflicts; namely it does the operation and if there is a conflict at verification time it falls back to two-phase locking. Another example is IndexFS’s inode cache which reduces RPCs by caching ancestor paths – the locality of this cache can be thrashed by random reads but performs well for metadata writes. For consistency, writes to directories in IndexFS block until the lease expires while writes to directories in ShardFS are slow for everyone as it either requires serialization or locking with many servers; reads in IndexFS are subject to cache locality while reads in ShardFS always resolve to 1 RPC. Another example of the overheads of addressing inconsistency is how CephFS maintains client sessions and inode caches for capabilities (which in turn make metadata access faster). When metadata is exchanged between metadata servers these sessions/caches must be flushed and new statistics exchanged with a scatter-gather process; this halts updates on the directories and blocks until the authoritative metadata server responds. These protocols are discussed in more detail in the next section but their inclusion here is a testament to the complexity of migrating metadata.

The conclusion we have drawn from this related work is that metadata protocols have a bigger impact on performance and scalability than load balancing. Understanding these protocols helps load balancing and gives us a better understanding of the metrics we should use to make migration decisions (*e.g.*, which operations reflect the state of the system), what types of requests cause the most load, and how an overloaded system reacts (*e.g.*, increasing latencies, lower throughput, etc.).

2.2 Relaxing POSIX

POSIX workloads require strong consistency and many file systems improve performance by reducing the number of remote calls per operation (*i.e.* RPC amplification). As discussed in the previous section, caching with leases and replication are popular approaches to reducing the overheads of path traversals but their performance is subject to cache locality and the amount of available resources, respectively; for random workloads larger than the cache extra RPCs hurt performance [?] and for write heavy workloads with more resources the RPCs for invalidations are harmful. Another approach to reducing RPCs is to use leases or capabilities.

High performance computing has unique requirements for file systems (*e.g.*, fast creates) and well-defined workloads (*e.g.*, workflows) that make relaxing POSIX sensible.

One popular approach is to allow clients to “lock” parts of the namespace to improve performance and scalability by avoiding synchronization, false sharing, and serialization. BatchFS assumes the application coordinates accesses to the namespace, so the clients can batch local operations and merge with a global namespace image lazily. Similarly, DeltaFS eliminates RPC traffic using subtree snapshots for non-conflicting workloads and middleware for conflicting workloads. MarFS gives administrators the ability to (1) lock “project directories” and (2) allocate GPFS clusters for demanding directory workloads. TwoTiers eliminates high-latencies by storing metadata in a flash tier; apps lock the namespace so that metadata can be accessed more quickly. Unfortunately, decoupling the namespaces has costs: (1) merging metadata state back into the global namespace is slow; (2) failures are local to the failing node; and (3) the systems are not backwards compatible.

For (1), state-of-the-art systems manage consistency in non-traditional ways: IndexFS maintains the global namespace but blocks operations from other clients until the first client drops the lease, BatchFS does operations on a snapshot of the namespace and merges batches of operations into the global namespace, and DeltaFS never merges back into the global namespace. The merging for BatchFS is done by an auxiliary metadata server running on the client and conflicts are resolved by the application. Although DeltaFS never explicitly merges, applications needing some degree of ground truth can either manage consistency themselves on a read or add a bolt-on service to manage the consistency.

For (2), if the client fails and stays down, all metadata operations on the decoupled namespace are lost. If the client recovers, the on-disk structures (for BatchFS and DeltaFS this is the SSTables used in TableFS) can be recovered. In other words, the clients have state that cannot be recovered if the node stays failed and any progress will be lost. This scenario is a disaster for checkpoint-restart where missed cycles may cause the checkpoint to bleed over into computation time.

For (3), decoupled namespace approaches sacrifice POSIX going as far as requiring the application to link against the systems they want to talk to. In today’s world of software defined caching, this can be a problem for large data centers with many types and tiers of storage. Despite well-known performance problems POSIX and REST are the dominant APIs for data transfer.

Decoupling the namespace delays metadata consistency and sacrifices fault tolerance. As shown in Table ??, metadata consistency is provided by capabilities and fault tolerance is addressed with a journal.

3. POSIX Overheads

In our examination of the overheads of POSIX we benchmark and analyze CephFS, the file system that uses the RADOS object store to store its data and metadata. We choose CephFS because it is an open-source production quality sys-

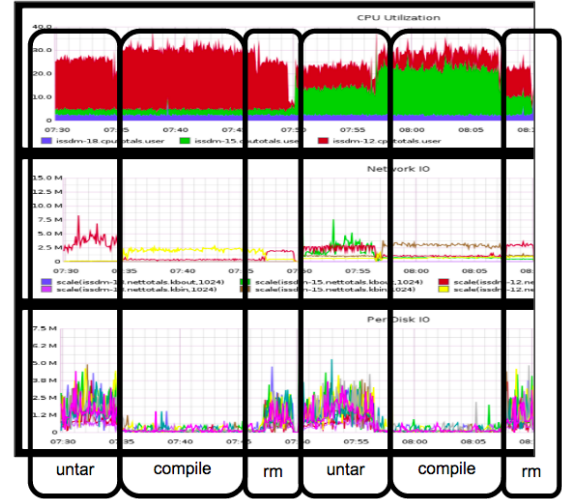


Figure 2. Create-heavy workloads (untar) incur the highest disk, network, and CPU utilization because the metadata server is managing consistency and fault tolerance.

tem. CephFS made one set of design decisions and we not asserting that the design decisions that were made are superior but instead highlight the effect those decisions have on performance.

To show how CephFS behaves under high metadata load we use a create-heavy workload because these types of workloads incur high CPU, network, and disk usage. Figure 2 shows a compilation of the Linux kernel, which has a download, untar, compile, and remove phase. The untar phase is characterized by many creates and has the highest resource usage, indicating that it is stressing the consistency and journalling subsystems of the metadata server. Also of note: a create-heavy workload does not help for caching inodes.

3.1 Fault Tolerance

As shown in Figure 3 CephFS uses RADOS (1) as a metadata store for all information about files including the hierarchical namespace and (2) as a staging area for the journal of updates before they are applied to the metadata store. (1) is essentially a cache that improves performance by serving metadata from memory and (2) is the mechanisms for achieving fault tolerance.

Fault tolerance means that the client or server can fail and metadata will not be lost. CephFS addresses fault tolerance with a metadata journal that streams into the resilient object store. Similar to LFS [1] and WAFL [2] the metadata journal can grow to large sizes ensuring (1) sequential writes into RADOS and (2) the ability for daemons to trim redundant or irrelevant journal entries.

The journal is striped over objects where multiple journal updates can reside on the same object. There are two tunables for controlling the journal: the segment size and the




Figure 3. CephFS uses a journal to stage updates and tracks dirty metadata in the collective memory of the MDSs. Each MDS maintains its own journal, which is broken up into 4MB segments. These segments are pushed into RADOS and deleted when that particular segment is trimmed from the end of the log. In addition to journal segments, RADOS also stores per-directory objects.




Figure 4. Performance improves with larger journal segments because the metadata server spends less time flushing the journal

number of parallel segments that can be written in parallel. The former is memory bound as larger segments take up more memory but can reduce the time spent journaling and the latter is CPU bound.

Figure 4 shows that journaling metadata updates into the object store has an overhead. Part (a) shows the runtime for different journal segment sizes; the larger the segment size the bigger that the writes into the object store are. The trade-off comes in terms of memory because larger segment sizes take up more space with their buffers. Parts (b), (c), and (d) show the throughput over time for different segment sizes. Performance suffers when time is spent journaling.

Despite this overhead, we posit that the journal is sufficient to slow down metadata throughput but not so much as to overwhelm RADOS because we measured our peak bandwidth to be 100MB/s, which is the speed of our network link.

Comparison to decoupled namespaces: In BatchFS and DeltaFS, as far as we can tell, when a client or server fails there is no recovery scheme. For BatchFS, if a client fails when it is writing to the local log-structured merged tree (implemented as an SSTable) then those batched metadata operations are lost. For DeltaFS, if the client fails then on

restart the computation does the work again – since the snapshots of the namespace are never globally consistent, there is no ground truth that requires the failed namespace to answer to anyone. On the server side, BatchFS and DeltaFS use IndexFS. Again, IndexFS writes metadata to SSTables but it is not clear whether they ever vacate memory, get written to disk, or are flushed to the object store.

3.2 Strong Consistency

Access to POSIX metadata is strongly consistent, so reads and writes are globally ordered. The synchronization and serialization machinery needed to ensure that all clients see the same state has high overhead. CephFS uses capabilities to keep metadata strongly consistent. To reduce the number of RPCs needed for consistency, clients can obtain capabilities for reading, reading and updating, reads caching, writing, buffering writes, changing the file size, and performing lazy IO.

To keep track of the read caching and write buffering capabilities, the clients and metadata servers agree on the state of each inode using an inode cache. If a client has the directory inode cached it can do metadata writes (e.g., create)




Figure 5. Journalling metadata updates has a bigger overhead than maintaining the inode cache. For create-heavy workloads the inode cache offers no performance benefits.

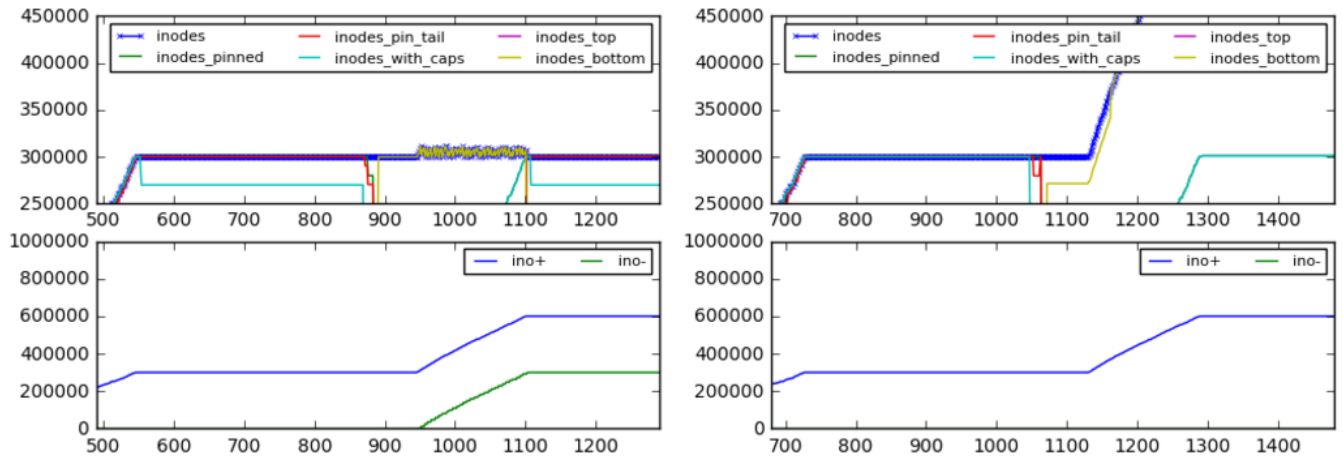


Figure 6. The inode cache improves metadata read performance but for our create-heavy workload it is only an overhead. Most of the time maintaining the cache is spent evicting and adding inodes.

with a single RPC. If the client is not caching the directory inode then it must do multiple RPCs to the metadata server to (1) determine if the file exists and (2) do the actual create. Unless the client immediately reads all the inodes in the cache, the inode cache is less useful for create-heavy workloads because the cached inodes are unused.

The benefits of caching the directory inode when creating files is shown in Figure 7(a). If only one client is creating files in a directory (“isolated” curve) then that client can lookup the existence of new files locally before issuing a create request to the metadata server. If another client starts creating files in the same directory (“interfere” curve) then the directory inode transitions out of read caching and the first client must send lookups to the metadata server. When other clients interfere the request throughput is higher Figure 7(b) but the runtime is slower because the isolated client scenario incurs less requests. Figures 7(c) and (d) plot the creates

and lookups, respectively, over time; creates slow down and lookups dominate the request load when the directory inode is shared in the interfering client scenario.

The drawbacks of the inode cache are shown in Figure ???. Figure ??(a) shows the throughput in metadata requests per second for a single client creating 200 thousand files. Until time 950 the throughput is steady at just under 2000 ops/sec and the CPU utilization is at about 30%. At time 950 seconds throughput degrades which corresponds to more inodes in the cache (Figure ??(c)). The values in Figure ??(c) are:

- inodes: total number of elemnts in the cache
- inodes pinned:
- inodes pinned tail
- inodes with caps
- inodes top




Figure 7. When a client create stream is “isolated” then lookups resolve locally but when a second client “interferes” by creating in the same directory, the directory inode capability is revoked forcing all clients to centralize lookups at the metadata server.

- inodes bottom

CephFS tries to keep the cache at 100 thousand inodes so the degradation in performance indicates that the metadata server cannot keep up with the workload. Figure ??(d) shows inodes getting added at the same rate as before 950 seconds (ino+) but the rate at which inodes get removed is less stable (ino-) indicating that cache eviction is taking more of the metadata servers time.

3.3 Consistency Overhead

Figure 8 is a baseline showing that the metadata server can service multiple clients when it is underloaded. One client creates files in the same directory and another does a touch or stat 15 seconds into the run. Figure 8(a) shows the runtime of the client doing the creates: “isolated” is when the create client is the only workload in the cluster, “interfere stat” is the runtime of the create client when another client does a stat, and “interfere touch” is the runtime of the create client when another client does a touch. There is only a minor performance degradation. Figure 8(b) compares the baseline the other client doing the touch or stat To explain Figure 8(a) we use Figure ??(b) to show why the runtime of creating 100 thousand files is unaffected. Compared to the throughput without an interfering stat (“isolated” curve) the throughput of the metadata server with an interfering stat (“interferring stat” curve) is higher. This means the metadata server is doing more operations suggesting that it can adequately handle the demands of both workloads.

To explain Figure 8(b) we use Figure 8(a) to show that operations are bounded by the number of files; more files incur more requests since both operations lookup every file. The number of files shown in Figure ??(a) indicates that the local stat (“client 0” bar) is interacting with less files than the remote stat (“client 1” bar). This graph also explains why the isolated stats and touches take the longest from the remote clients; out of all setups, the isolated operations are

doing the most RPCs. The difference between local (“client 0 interfere” bar) and remote (“client 1 interfere” bar) in Figure 8(b) is the overhead of doing RPCs. Isolated is the slowest because it requests the most files. On average the local interfering client (“client 0” bar in Figure ??(a)) requests 40 thousand less files than the remote interfering call (“client 1” bar in Figure ??(a)). This reflects how fast the local client can get into the queue of requests, presumably because it bypasses capability checks.

From this experiment we make 2 conclusions:

- the metadata server is not overloaded because it can handle both workloads
- the metadata server needs to get global consistency from only 1 client

Comparison to decoupled namespaces: Decoupled namespaces merge batches of metadata operations into the global namespaces when the job completes. In BatchFS the merge is delayed by the application using an API to switch between asynchronous to synchronous mode. The merge itself is explicitly managed by the application but future work looks at more automated methodologies. In DeltaFS snapshots of the metadata subtrees stays on the client machines; there is no ground truth and consistent namespaces are constructed and resolved at application read time or when a 3rd party system (e.g., middleware, scheduler, etc.) needs a view of the metadata.

4. Methodology

4.1 Decoupled Namespaces in CephFS

Recall that CephFS uses RADOS (1) as a metadata store for all information about files including the hierarchical namespace and (2) as a staging area for the journal of updates before they are applied to the metadata store. The journal tool is used for disaster recovery and lets administrators view and modify the journal of metadata updates; it can read the jour-

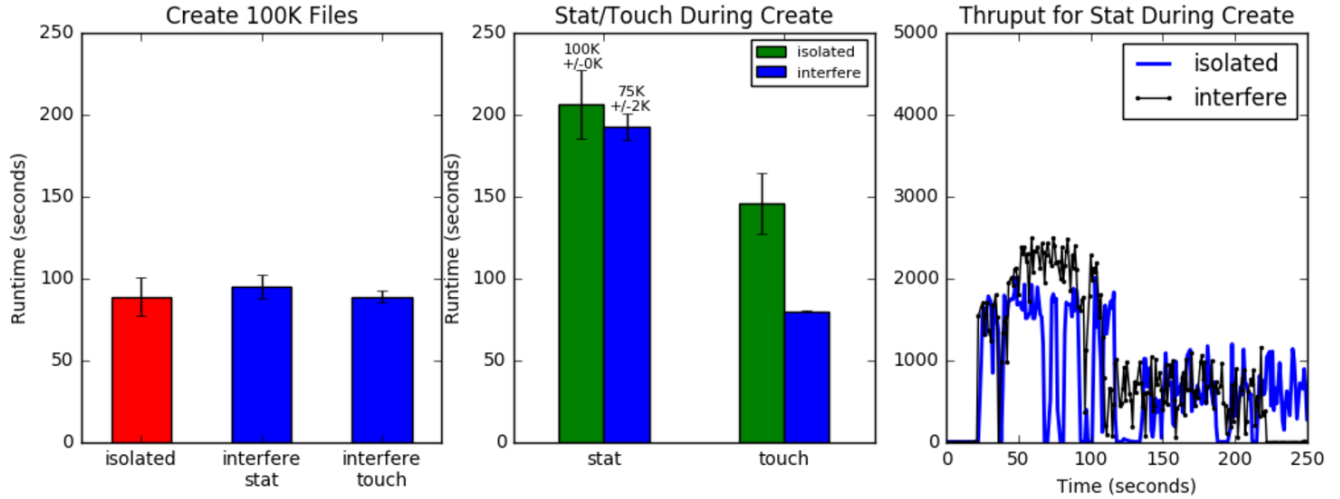


Figure 8. An underloaded metadata server adequately services conflicting clients; the create speeds are similar while the interfering operations are limited by the cost of RPCs.




Figure 9. Scaling clients shows increased variability when another client interferes; zooming in on runs with 7 clients we see that different types of interfering operations have different effects on performance variability and predictability.




Figure 10. Applications that can tolerate weaker consistency can decouple the namespace, write updates to a local file, and delay metadata updates. The merge and replay steps are optional if there is no need for a global namespace.

nal, erase events from the journal, and apply the updates in the journal to the metadata store. To apply journal updates to the metadata store, the journal tool reads the journal seg-

ments from RADOS objects and applies the update to the metadata store (which are also stored as RADOS objects).

The journal segments are saved as objects in RADOS. The journal has 4 pointers, described in 'osdc/Journaler.h':

- write position: tail of the journal; points to the current session where we are appending events
- unused field: where someone is reading
- expire position: old journal segments
- trimmed position: where daemon is expiring old items

Journal segments in RADDS have a header followed by serialized log events. The log events are read by hopping over objects using the read offset and object size pulled from the journal header. After decoding them, we can examine the metadata (1) about the event (e.g., type, timestamps, etc.) and (2) for inodes that the event touches.

The metadata for inodes that the event touches are called metadata blobs and the ones associated with events are **unordered**; this layout makes writing journal updates fast but the cost is realized when reading the metadata blobs. It makes sense to optimize for writing since reading only occurs on failures. To reconstruct the namespace for the metadata blob, the journal tool iterates over each metadata blob in the events and builds mappings of inodes to directory entries (for directories) and parent inodes to child inodes/directory entries.

The journal tool imports journals from binary files stored on disk. First the header of the dump is sanity checked and written into RADOS to the “header” object. The “header” object has metadata about the journal as well as the locations of all the journal pointers (e.g., where the tail of the journal is, where we are currently trimming, etc.). Then the journal events are cleaned (erasing trailing events that are not part of the header) and written as objects into RADOS. Note that while the journal is in RADOS, the metadata servers do not have the namespace reconstructed in memory so the metadata cluster will not service requests relating to the journal of imported events. To construct the namespace in the collective memory of the metadata servers we need to first construct the namespace in RADOS. The journal tool can explicitly do this by applying the journal to the metadata store in RADOS. This will pull the objects containing journal segments and replay them on the metadata store. Finally, we delete the journal in RADOS and restart the metadata servers so they rebuild their caches.

The journal tool exports journals to binary files stored on disk. First the journal is scanned for the header and then journal is recovered. To recover the journal the “header” object is read off disk and then objects are probed in order and starting from the write position saved in the header. Probing will update the write position if it finds objects with data in them. There are two types of objects:

1. journal objects (e.g., 200.000*)
2. metadata store objects (e.g., 1.inode)

We need to update the journal objects with the new events! How do we write out new objects to RADOS (e.g., 200.000*)? These should be the correct sizes (e.g., 4980 bytes) – these log events that we are adding are not making them into the journal objects... do we have to explicitly read them and then re-write them?

TODO: pull the object, read it, append to it, make sure size changes. It looks like we can only read these objects and apply doesn’t actually touch them... it only touches the metadata store objects (e.g., 1.inode*).

When exporting a journal of events, the journal tool first scans the journal to check for corruption. Then it recovers the journal by reading the “header” object out of RADOS. After reading the header, the journal tool can pull journal segments from RADOS because it knows how many objects to pull and how far to seek within those objects.

4.1.1 Leveraging the Journal Tool

We leverage the journal tool’s ability to reconstruct the namespace for the metadata blob. First we scan journal and store all events in a map. Next we create a new event and copy the root inode directory lump from another journal event into the event’s metadata blob. Finally, we add a new directory lump for our new file or directory.

If decoupled namespaces are turned on, the client writes metadata updates locally and merges the updates with the journal tool. The client that decouples the namespace operates without any consistency and any conflicts at the merge are resolved in favor of this client. Updates by other clients (*i.e.* metadata writes to the global namespace) are overwritten. Cudele adopts the following process when an application decouples the namespace:

1. metadata server exports a snapshot of the file system to a binary file
2. client materializes the snapshot in memory
3. client operates on the data structures in memory
4. client exports the new snapshot to a binary file
5. metadata server merges the snapshot into the global namespace

Step 3 is the most complicated and requires understanding how the snapshot is materialized in memory.

4.1.2 Operating on Snapshots

Our first implementation attempted to re-create journal events using the same libraries that the metadata server uses. To construct a `mkdir` we tried to instantiate a Ceph inode and directory entry for the current file/dir and its parent. This is too hard because there are too many moving parts in the metadata server (e.g., a `mdlog` class, stuff in memory, assumption that we can traverse up and down namespace, etc.). So when I tried add dentries and inodes it was trying to traverse up/down and it would almost always segfault when it was looking for something. These metablobs are supposed to be self container – the problem is I do not know what is supposed to go **inside** them.

Our second idea was to copy the metadata blob and change just what we needed. For example, we would save a binary dump of a generic `mkdir` event on disk. When the application makes a directory, this dump would be loaded and the fields would be changed before being written back to disk. Rather than traversing up and down a namespace in memory of a metadata server, we should traverse up and down the namespace **inside** the metadata blob. This implementation requires disk IO and editing the log event is non-trivial for two reasons:

- methods do not edit events; they just write them
- the metadata that the event touches (e.g., the metablob) is unorganized on disk for performance – it is trade-off

for writing data faster serially and reconstructing information slowly since failure is not considered the norm

Faced with these challenges we landed on our final implementation: load the snapshot into the the data structures used to examine and replay journals, edit those data structures, and write them out to disk as binary.

5. Results

In this section, we quantify the costs of consistency and fault tolerance in CephFS. If the components that ensure these semantics (i.e. capabilities and journals, respectively) can be mitigated or delayed, then global namespaces perform as well as decoupled namespaces.

Experimental Setup: we run our experiments on a 9 OSD, 3 MDS, 1 MON Ceph cluster. The clients use the Ceph kernel client, which has been in the mainline Linux kernel since TODO. We use the kernel client so that we can find the true create speed of the server; our experiments show a low CPU utilization for the clients which indicates that we are stressing the servers more. We also turn caching off because, as shown in figureX there is little difference, in terms of performance between caching and not caching when using the kernel client.

5.1 Journaling Overhead

Journal to RADOS Turn off journaling (large segment) Journal to in-memory OSD

5.2 Macrobenchmarks

updatedb: <http://lists.ceph.com/pipermail/ceph-users-ceph.com/2015-July/002768.html>

References

- [1] Christina L. Abad, Huong Luu, Yi Lu, and R Campbell. Metadata Workloads for Testing Big Storage Systems. Technical report, Citeseer, 2012.
- [2] Cristina L. Abad, Huong Luu, Nathan Roberts, Kihwal Lee, Yi Lu, and Roy H. Campbell. Metadata Traces and Workload Models for Evaluating Big Storage Systems. In *Proceedings of the 2012 IEEE/ACM Fifth International Conference on Utility and Cloud Computing*, UCC '12, pages 125–132, 2012.
- [3] Sadaf R. Alam, Hussein N. El-Harake, Kristopher Howard, Neil Stringfellow, and Fabio Verzelloni. Parallel I/O and the Metadata Wall. In *Proceedings of the 6th Workshop on Parallel Data Storage*, PDSW'11, 2011.
- [4] John Bent, Brad Settlemeyer, and Gary Grider. Serving Data to the Lunatic Fringe.
- [5] Kirk McKusick and Sean Quinlan. GFS: Evolution on Fast-forward. *Communications ACM*, 53(3):42–49, March 2010.
- [6] Kai Ren, Qing Zheng, Swapnil Patil, and Garth Gibson. IndexFS: Scaling File System Metadata Performance with Stateless Caching and Bulk Insertion. In *Proceedings of the 20th ACM/IEEE Conference on Supercomputing*, SC '14, 2014.
- [7] Drew Roselli, Jacob R. Lorch, and Thomas E. Anderson. A Comparison of File System Workloads. In *Proceedings of the Annual Conference on USENIX Annual Technical Conference*, ATEC '00, pages 4–4, 2000.
- [8] Michael A. Sevilla, Noah Watkins, Carlos Maltzahn, Ike Nassi, Scott A. Brandt, Sage A. Weil, Greg Farnum, and Sam Fineberg. Mantle: A programmable metadata load balancer for the ceph file system. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, SC '15, 2015.
- [9] Ashish Thusoo, Zheng Shao, Suresh Anthony, Dhruba Borthakur, Namit Jain, Joydeep Sen Sarma, Raghotham Murthy, and Hao Liu. Data Warehousing and Analytics Infrastructure at Facebook. In *Proceedings of the 2010 ACM SIGMOD International Conference on Management of Data*, SIGMOD '10, 2010.
- [10] Sage A. Weil, Scott A. Brandt, Ethan L. Miller, Darrell D. E. Long, and Carlos Maltzahn. Ceph: A Scalable, High-Performance Distributed File System. In *Proceedings of the 7th USENIX Symposium on Operating Systems Design & Implementation*, OSDI'06, 2006.